

Drilling Security Data

Charles S. Givre, CISSP
charles.givre@gtkcyber.com
www.thedataist.com



2 Huge Problems in Data Analytics

- Data is often stored in multiple disparate systems, and querying these systems requires understanding multiple query languages
- Joining data from these multiple disparate systems is difficult, costly and time consuming

Que feriez-vous si toutes vos données parlaient la même langue?

Què faria vostè si totes les seves dades parlaven el mateix idioma?

ሁሉም የእርስዎ ውሂብ ተመሳሳይ ቋንቋ ሲጠቀም ምን ያደርጉ ነበር?

What would you do if all your data spoke the same language?

Что бы вы сделали, если все ваши данные говорили на одном языке?

מה היית עושה אם כל הנתונים שלך מדברים באותה שפה

اگر تمام داده های شما به همان زبان صحبت می کنند چه کاری انجام می دهید؟

Was würden Sie tun, wenn alle Ihre Daten die gleiche Sprache sprechen würden?

The Problem: Analyzing Security Data is Hard

Security Data Analysis is Hard...Really Hard



Why?

Security Data comes in Many Forms

Security Data Comes in Many Forms

- Standard Types such as JSON/CSV/XML
- Log Files: Event Logs, Database, Web Server etc.
- Syslog
- PCAP / PCAP-NG: Binary, raw network traffic

Security Data Lives In Many Places

- File systems
- Cloud Storage: Azure / S3
- Databases
- Real Time Event Streams
- Other sources

Few tools can effectively analyze
these data types effectively

Even fewer tools can effectively
analyze **ALL** these data types
effectively

Splunk and ELK are solid
SIEM platforms

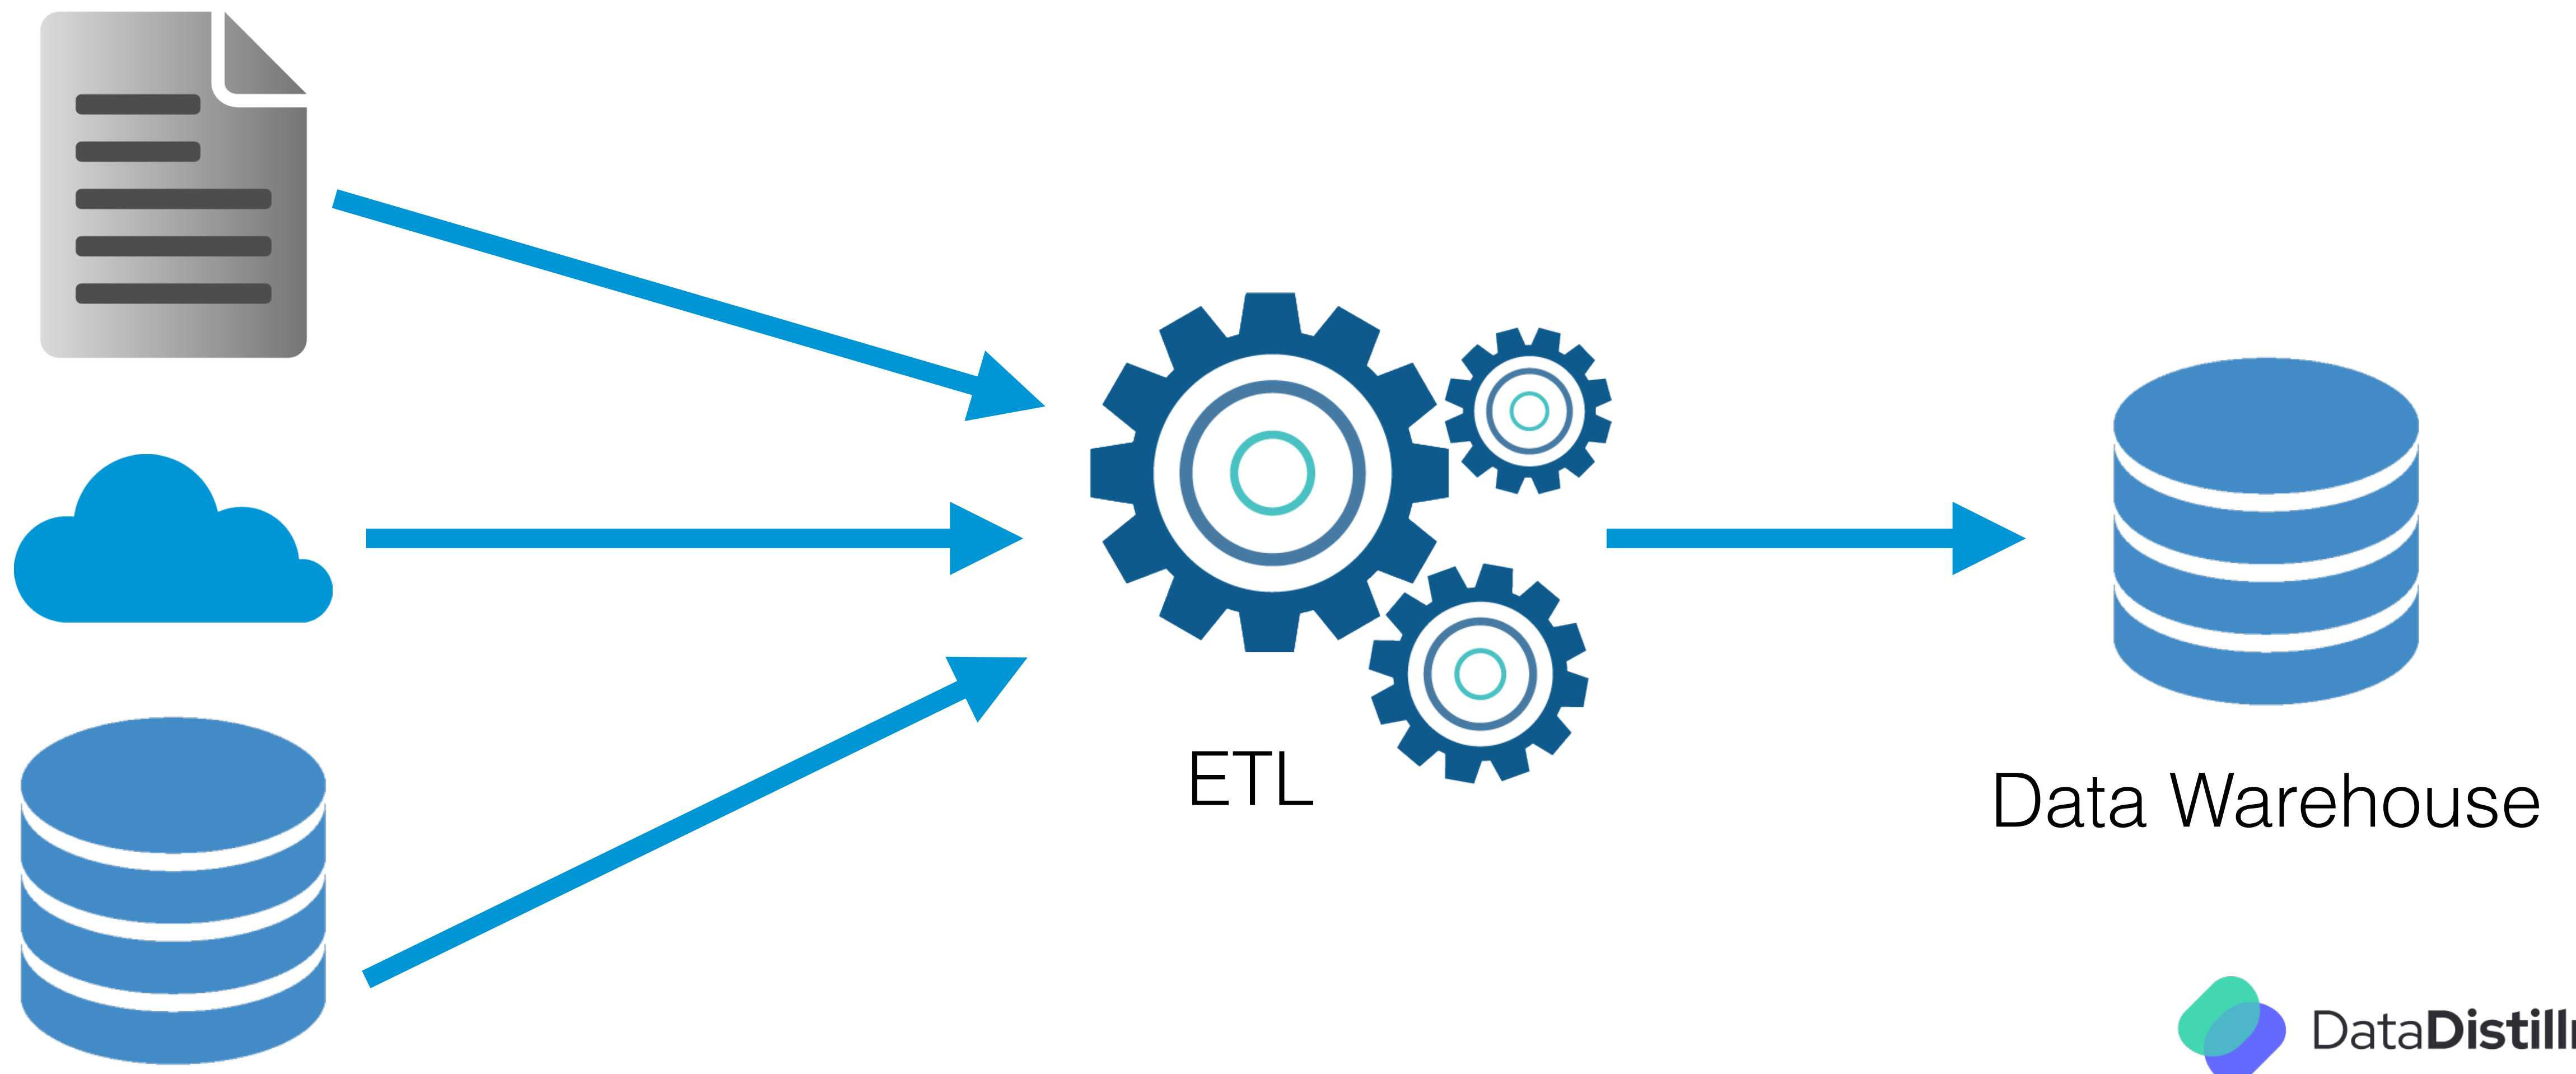


Wireshark is good for PCAP
analysis



Security data is not arranged in
an optimal way for ad-hoc
analysis

Security data is not arranged in an optimal way for ad-hoc analysis



ETL is expensive and
wasteful

Analytics teams spend between 50%-90% of their time preparing their data.

76% of Data Scientists say this is
the **least enjoyable part** of their
job.

The ETL Process **consumes the most time** and **contributes almost no value** to the end product.

8 Wastes

The 8 Wastes are eight types of process obstacles that get in the way of providing value to the customer.



Defects

Efforts caused by rework, scrap, and incorrect information.



Overproduction

Production that is more than needed or before it is needed.



Waiting

Wasted time waiting for the next step in a process.



Non-Utilized Talent

Underutilizing people's talents, skills, & knowledge.



Transportation

Unnecessary movements of products & materials.



Inventory

Excess products and materials not being processed.



Motion

Unnecessary movements by people (e.g., walking).



Extra-Processing

More work or higher quality than is required by the customer.



So where does Drill fit in?



Apache Drill is an SQL Engine
for self-describing data.






Drill lets you query anything*,
wherever it is*, no matter its
size** using standard SQL.



* well.. almost anything
** within reason

Drill acts as a universal translator for data.



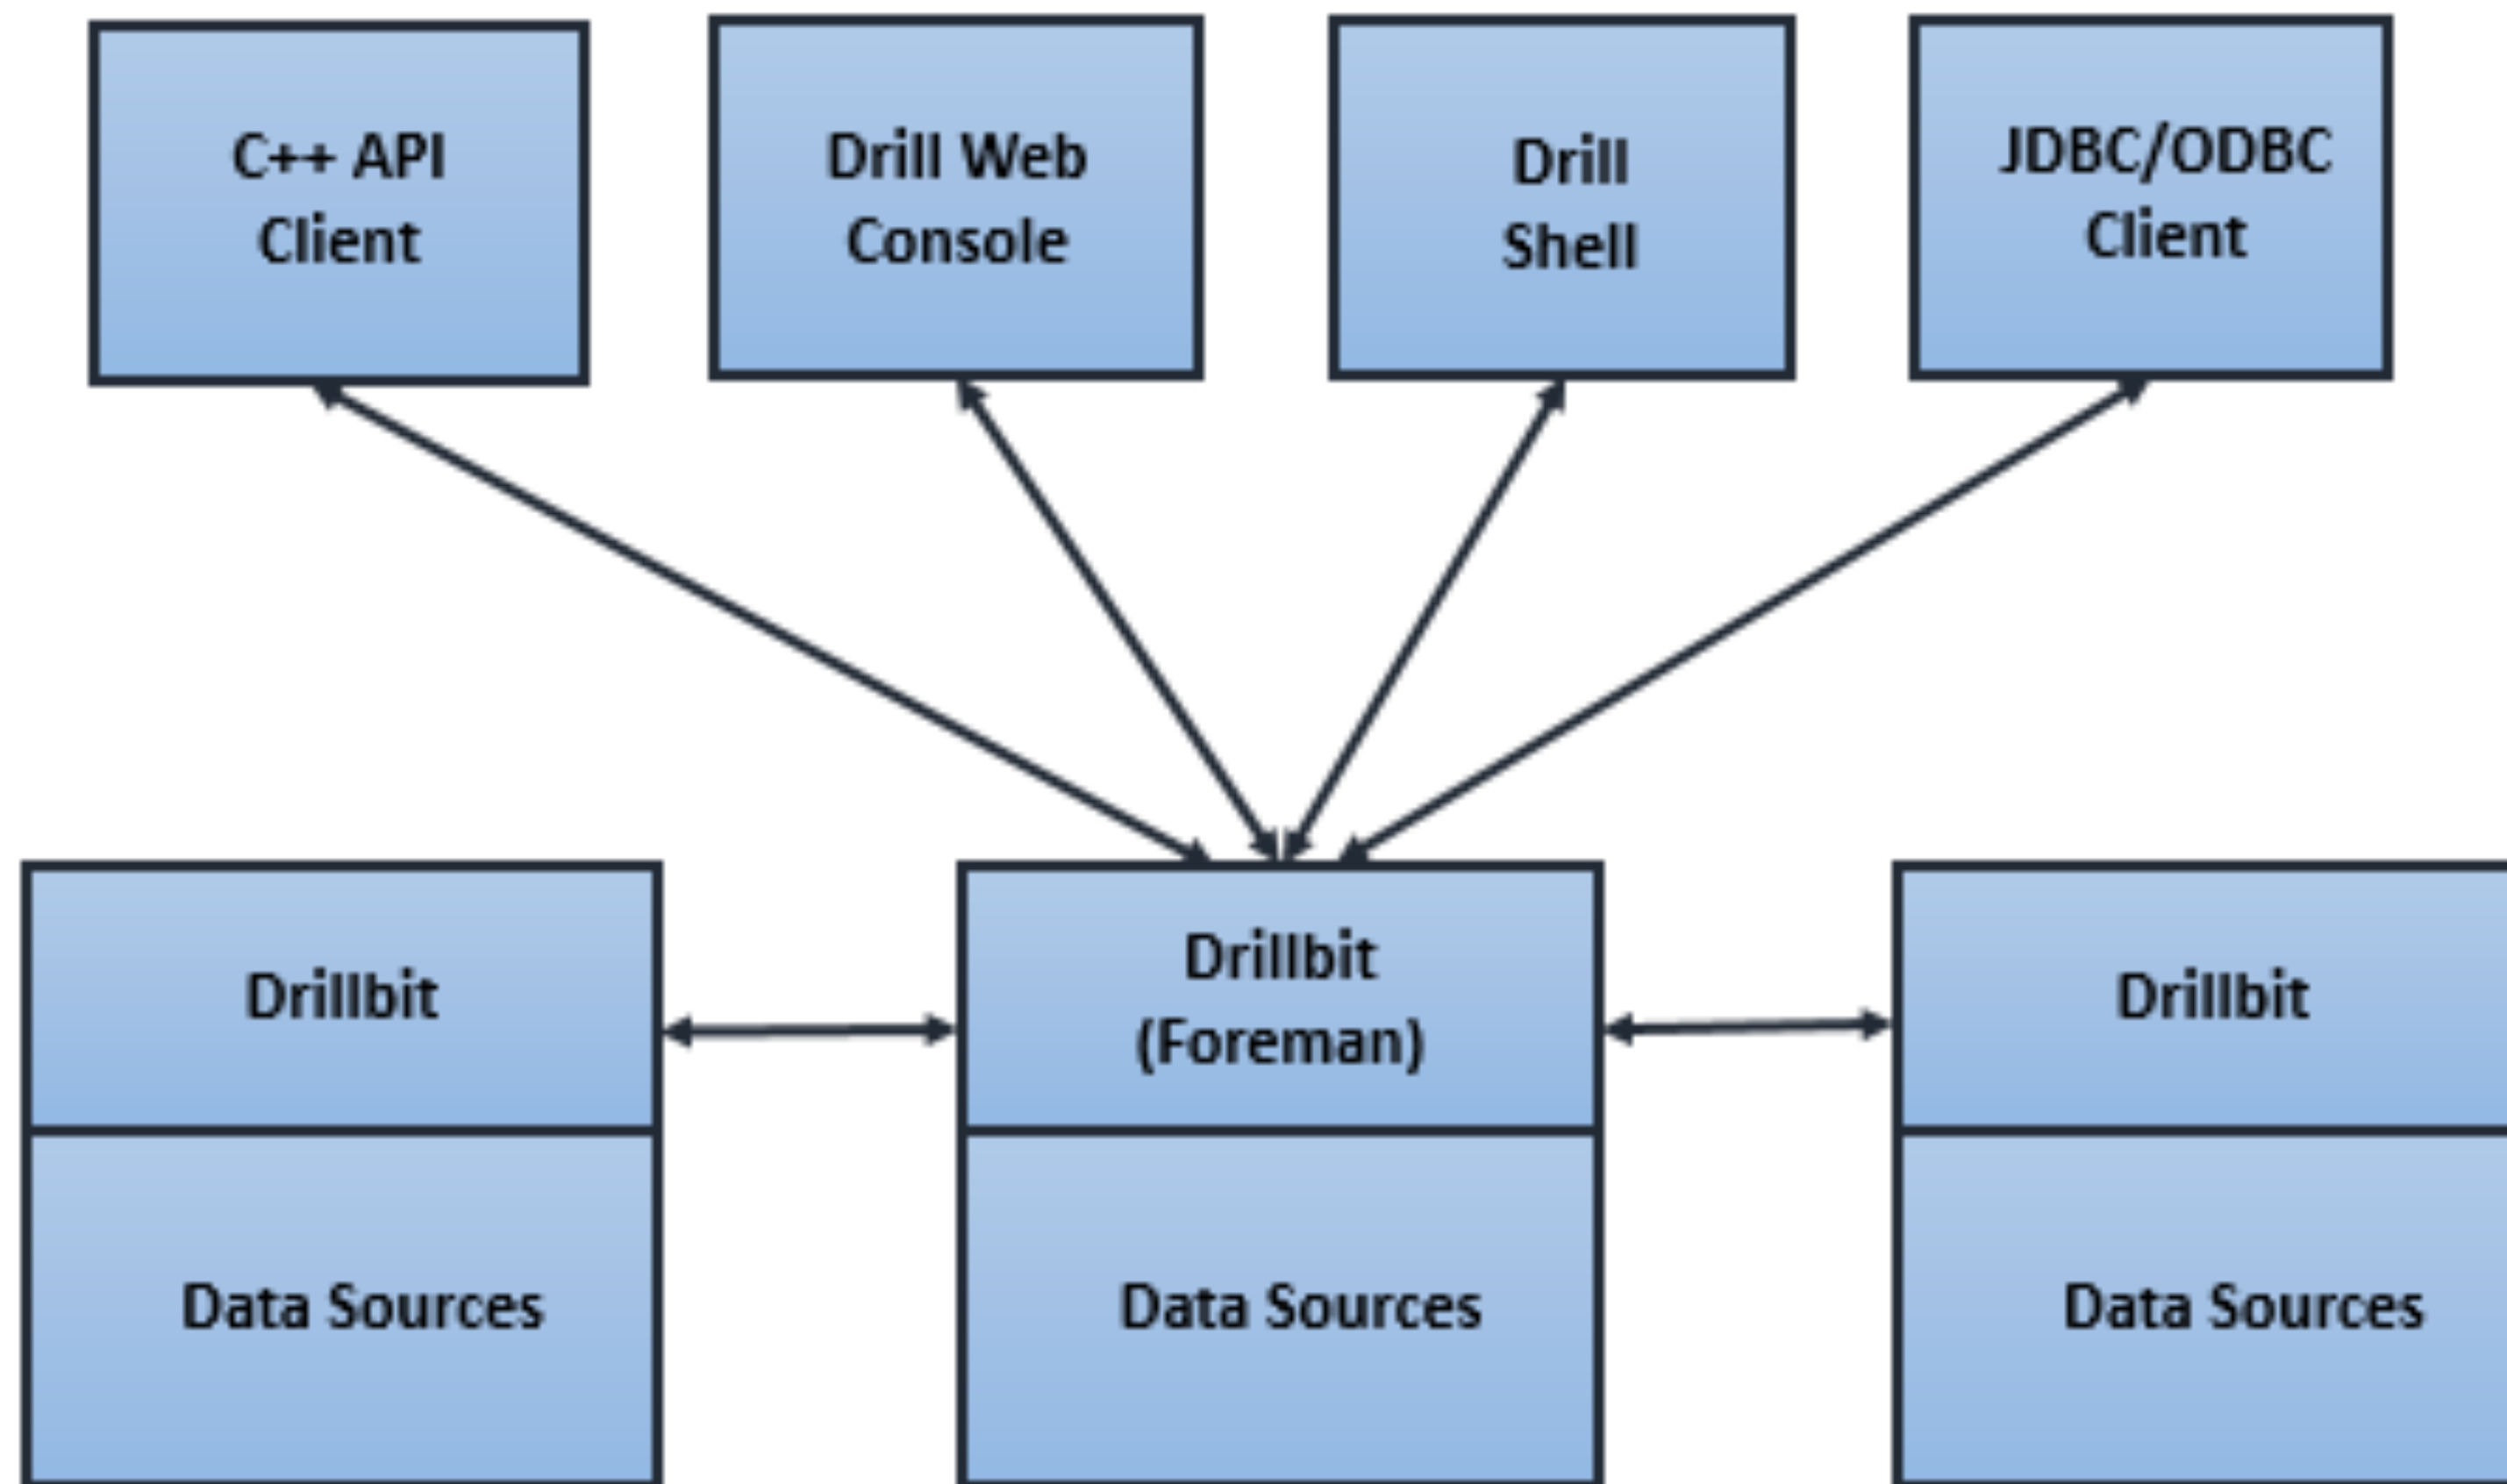







How Drill Works

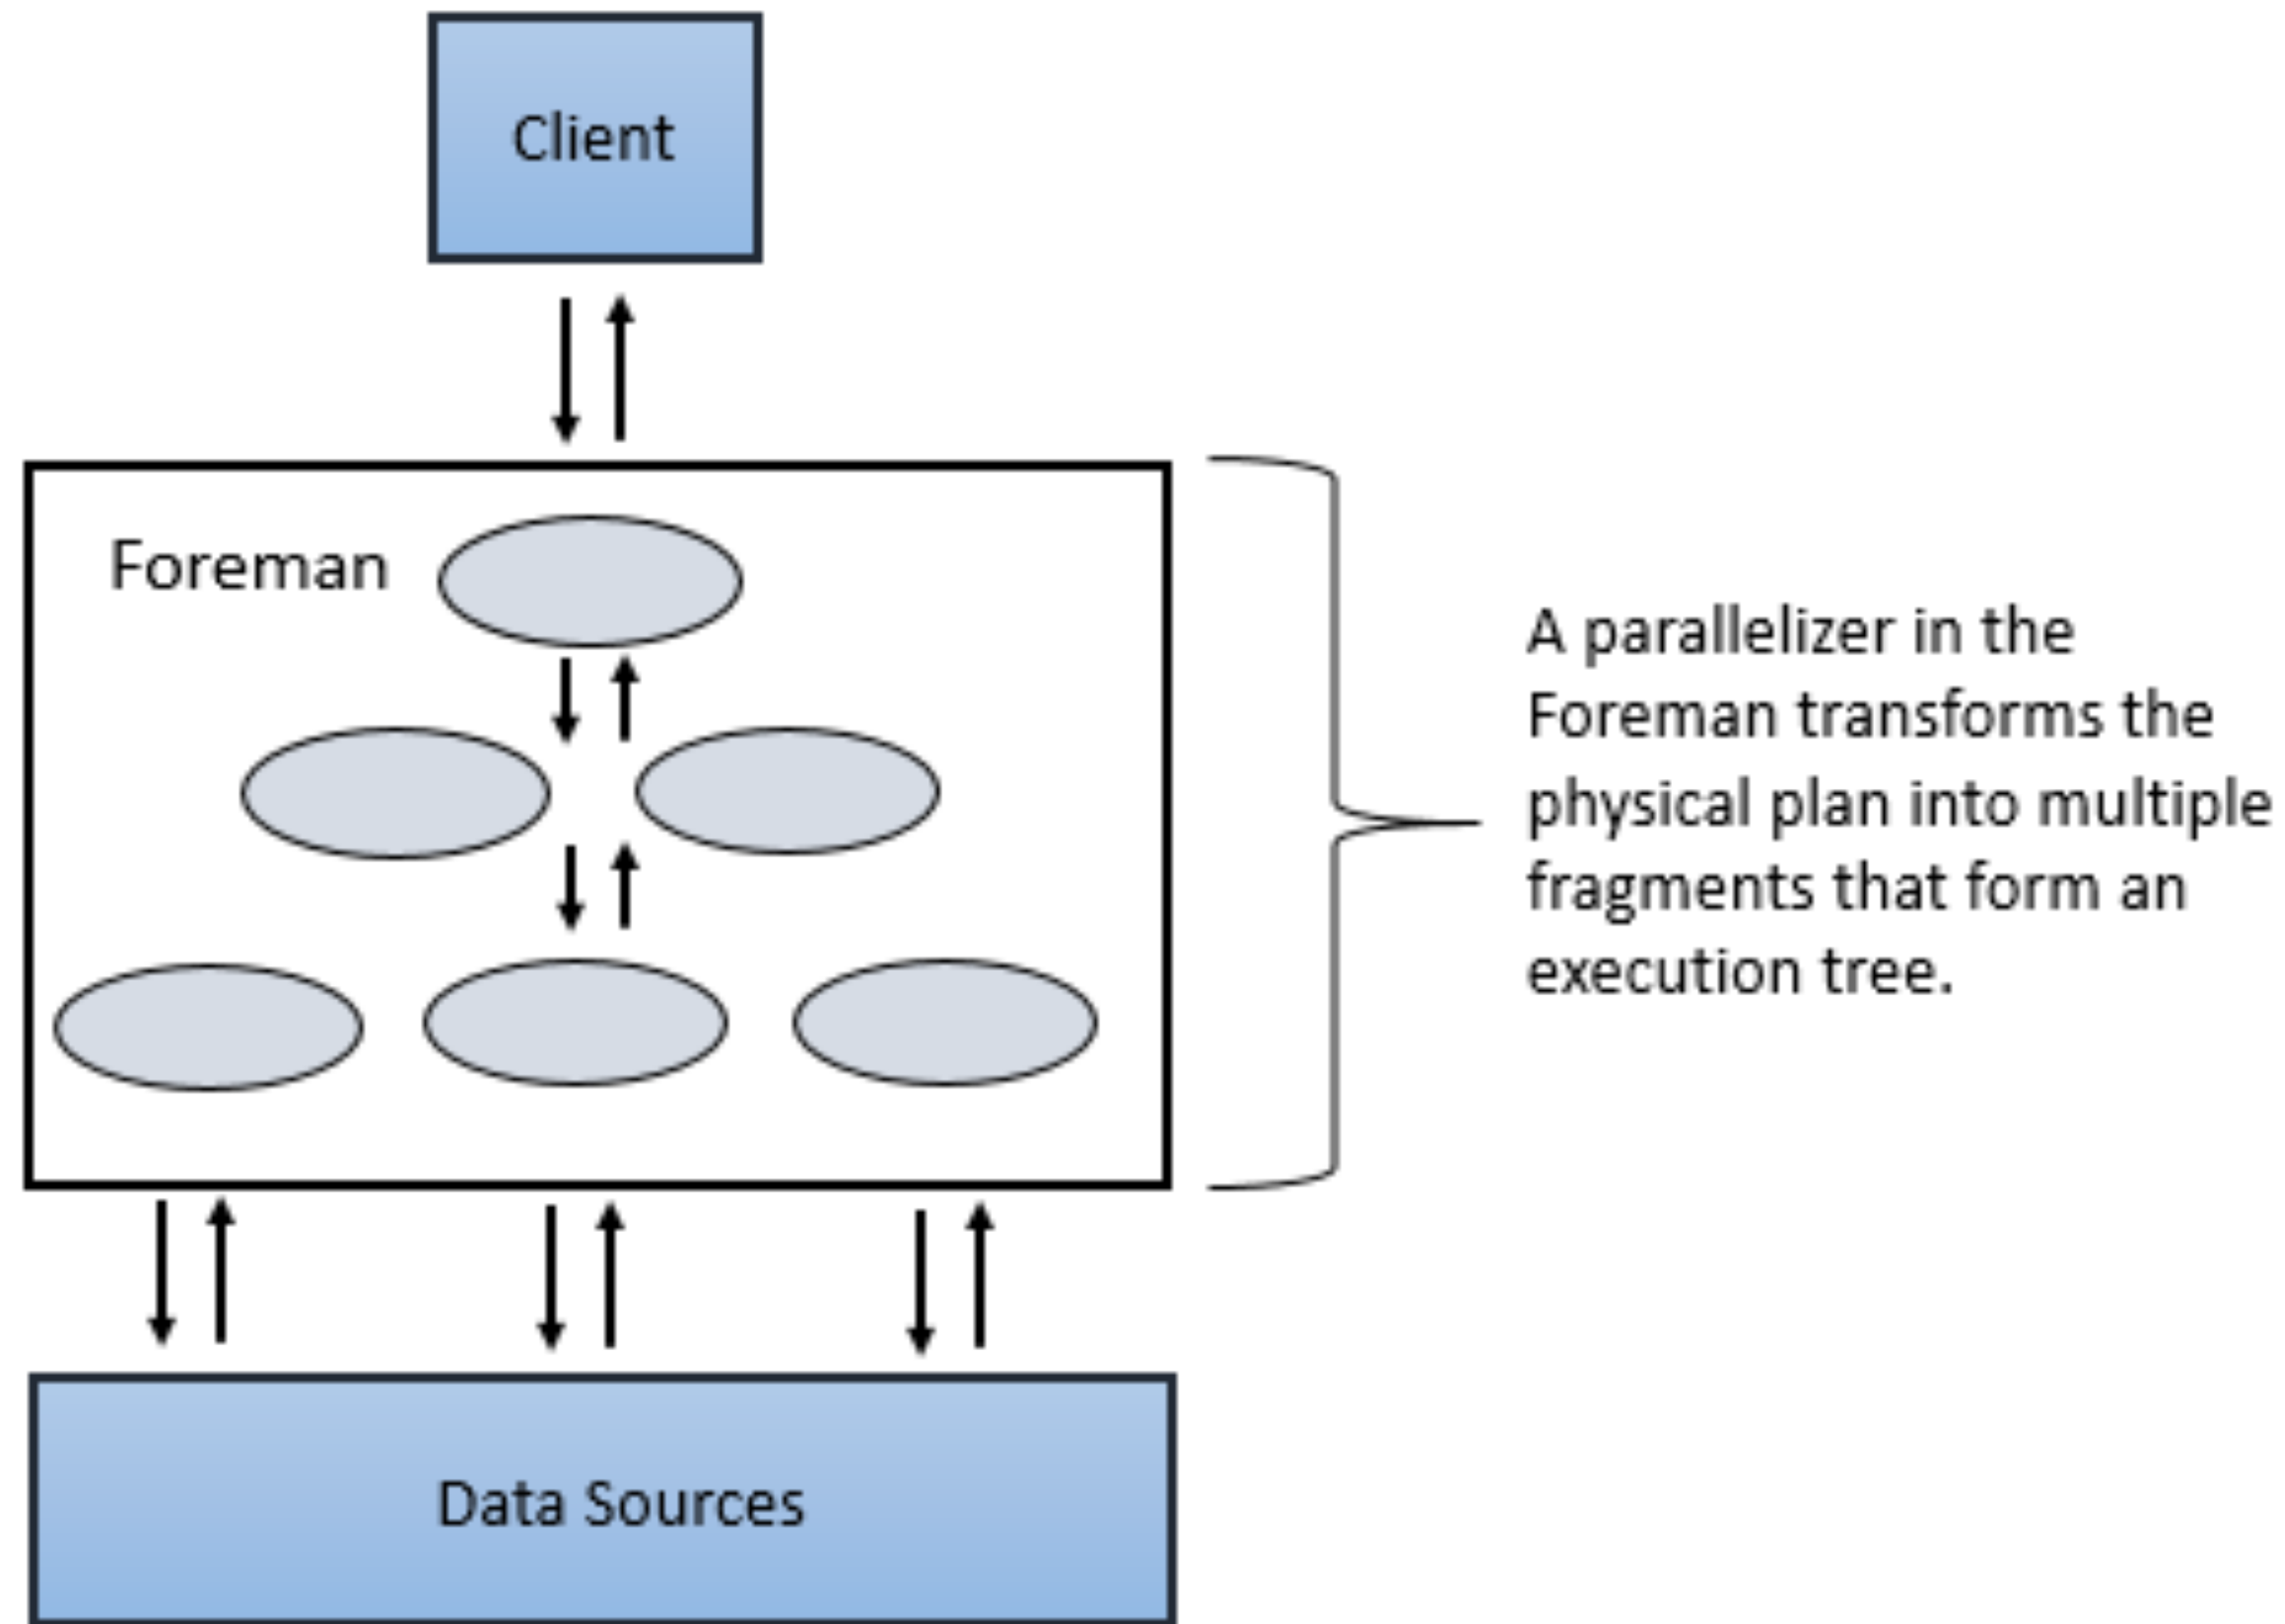
How Drill Works



How Drill Works



How Drill Works



So let's take a look...

```
158.222.5.157 - - [25/Oct/2015:04:24:37 +0100] "GET /acl_users/  
credentials_cookie_auth/require_login?came_from=http%3A//howto.basjes.nl/  
join_form HTTP/1.1" 200 10716 "http://howto.basjes.nl/join_form" "Mozilla/5.0  
(Windows NT 6.3; WOW64; rv:34.0) Gecko/20100101 Firefox/34.0 AlexaToolbar/  
alxf-2.21"
```

```
158.222.5.157 - - [25/Oct/2015:04:24:39 +0100] "GET /login_form HTTP/1.1" 200  
10543 "http://howto.basjes.nl/" "Mozilla/5.0 (Windows NT 6.3; WOW64; rv:34.0)  
Gecko/20100101 Firefox/34.0 AlexaToolbar/alxf-2.21"
```

📄 `hackers-access.httpd` 🔍 📄 ✕

request_referer_ref	VARCHAR
request_receive_time_last_time	VARCHAR
request_firstline_uri_protocol	VARCHAR
request_receive_time_microsecond	BIGINT
request_receive_time_last_microsecond__utc	BIGINT
request_firstline_original_uri_query_\$	UserDefinedType
request_firstline_original_protocol	VARCHAR
request_firstline_original_uri_host	VARCHAR
request_referer_host	VARCHAR
request_receive_time_month__utc	BIGINT
request_receive_time_last_minute	BIGINT
request_firstline_protocol_version	VARCHAR
request_receive_time_time__utc	VARCHAR
request_referer_last_ref	VARCHAR
request_receive_time_last_timezone	VARCHAR
request_receive_time_last_weekofweekyear	BIGINT
request_referer_last	VARCHAR
request_receive_time_minute	BIGINT
connection_client_host_last	VARCHAR
request_receive_time_last_millisecond__utc	BIGINT
request_firstline_original_uri	VARCHAR
request_firstline	VARCHAR
request_receive_time_nanosecond	BIGINT
request_receive_time_last_millisecond	BIGINT
request_receive_time_day	BIGINT
request_referer_port	BIGINT
request_firstline_original_uri_port	BIGINT
request_receive_time_year	BIGINT
request_receive_time_last_date	VARCHAR
request_referer_query_\$	UserDefinedType

request_referer_last	request_receive_time_minute	connection_client_host_last	request_receive_time_last_millisecond__utcrequest_firstline_original_uri	request_firstline_original_uri
http://howto.basjes.nl/join_form				
http://howto.basjes.nl/	11	195.154.46.135	0	/linux/doing-pxe-without-dhcp-control
http://howto.basjes.nl/	11	23.95.237.180	0	/join_form
http://howto.basjes.nl/join_form	11	23.95.237.180	0	/join_form
http://howto.basjes.nl/	24	158.222.5.157	0	/join_form
http://howto.basjes.nl/join_form	24	158.222.5.157	0	/join_form
http://howto.basjes.nl/join_form	24	158.222.5.157	0	/acl_users/credentials_cookie_auth/require_login?came_from=http%
http://howto.basjes.nl/	24	158.222.5.157	0	/login_form
http://howto.basjes.nl/login_form	24	158.222.5.157	0	/login_form
http://howto.basjes.nl/	32	5.39.5.5	0	/join_form
http://howto.basjes.nl/	34	180.180.64.16	0	/linux/doing-pxe-without-dhcp-control
http://howto.basjes.nl/	34	180.180.64.16	0	/join_form
http://howto.basjes.nl/join_form	34	180.180.64.16	0	/join_form
http://howto.basjes.nl/join_form	34	180.180.64.16	0	/acl_users/credentials_cookie_auth/require_login?came_from=http%
http://howto.basjes.nl/	34	180.180.64.16	0	/login_form

Who is Trying to Hack This Site?

- The url /join_form is not public so anyone attempting to access this site, so almost anyone trying to access this probably a hacker...
- Let's see who is looking...

Who is Trying to Hack This Site?

```
SELECT request_receive_time,  
connection_client_host,  
request_useragent  
FROM dfs.test.`hackers-access.httpd`  
WHERE request_firstline_uri = '/join_form'
```

Who is Trying to Hack This Site?

```
SELECT request_receive_time,  
connection_client_host,  
request_useragent  
FROM dfs.test.`hackers-access.httpd`  
WHERE request_firstline_uri = '/join_form'
```

Results

Query History

Preview: `hackers-access.httpd`

237.180

Explore


.CSV

Clipboard

Search Results

request_receive_time	connection_client_host	request_useragent
2015-10-25T03:11:26	23.95.237.180	Mozilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20100101 Firefox/35.0
2015-10-25T03:11:27	23.95.237.180	Mozilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20100101 Firefox/35.0
2015-10-25T03:24:31	158.222.5.157	Mozilla/5.0 (Windows NT 6.3; WOW64; rv:34.0) Gecko/20100101 Firefox/34.0 AlexaToolbar/alxf-2.21
2015-10-25T03:24:32	158.222.5.157	Mozilla/5.0 (Windows NT 6.3; WOW64; rv:34.0) Gecko/20100101 Firefox/34.0 AlexaToolbar/alxf-2.21
2015-10-25T03:32:22	5.39.5.5	Mozilla/5.0 (Windows NT 5.1; rv:34.0) Gecko/20100101 Firefox/34.0
2015-10-25T03:34:40	180.180.64.16	Mozilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20100101 Firefox/35.0
2015-10-25T03:34:42	180.180.64.16	Mozilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20100101 Firefox/35.0
2015-10-25T04:06:42	89.42.237.71	Mozilla/5.0 (Windows NT 5.1; rv:34.0) Gecko/20100101 Firefox/34.0
2015-10-25T04:06:43	89.42.237.71	Mozilla/5.0 (Windows NT 5.1; rv:34.0) Gecko/20100101 Firefox/34.0

Who is Trying to Hack This Site?



Let's take a look at those IP
addresses shall we?

Who is Trying to Hack This Site?

```
SELECT request_receive_time,  
connection_client_host,  
request_useragent  
FROM dfs.test.`hackers-access.httpd`  
WHERE request_firstline_uri = '/join_form'
```

Results

Query History

Preview: `hackers-access.httpd`

237.180

Explore

.CSV

Clipboard

Search Results

request_receive_time	connection_client_host	request_useragent
2015-10-25T03:11:26	23.95.237.180	Mozilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20100101 Firefox/35.0
2015-10-25T03:11:27	23.95.237.180	Mozilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20100101 Firefox/35.0
2015-10-25T03:24:31	158.222.5.157	Mozilla/5.0 (Windows NT 6.3; WOW64; rv:34.0) Gecko/20100101 Firefox/34.0 AlexaToolbar/alxf-2.21
2015-10-25T03:24:32	158.222.5.157	Mozilla/5.0 (Windows NT 6.3; WOW64; rv:34.0) Gecko/20100101 Firefox/34.0 AlexaToolbar/alxf-2.21
2015-10-25T03:32:22	5.39.5.5	Mozilla/5.0 (Windows NT 5.1; rv:34.0) Gecko/20100101 Firefox/34.0
2015-10-25T03:34:40	180.180.64.16	Mozilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20100101 Firefox/35.0
2015-10-25T03:34:42	180.180.64.16	Mozilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20100101 Firefox/35.0
2015-10-25T04:06:42	89.42.237.71	Mozilla/5.0 (Windows NT 5.1; rv:34.0) Gecko/20100101 Firefox/34.0
2015-10-25T04:06:43	89.42.237.71	Mozilla/5.0 (Windows NT 5.1; rv:34.0) Gecko/20100101 Firefox/34.0

Drill's Flexible UDF Interface
Allows you to write your own
functions.

A collection of GeolP Functions
is available on GitHub.

<https://github.com/cgivre/drill-geoip-functions>

Drill GeoIP Functions

- **getCountryName(<ip>)**: This function returns the country name of the IP address, "Unknown" if the IP is unknown or invalid.
- **getCountryConfidence(<ip>)**: This function returns the confidence score of the country ISO code of the IP address.
- **getCountryISOCode(<ip>)**: This function returns the country ISO code of the IP address, "Unknown" if the IP is unknown or invalid.
- **getCityName(<ip>)**: This function returns the city name of the IP address, "Unknown" if the IP is unknown or invalid.
- **getCityConfidence(<ip>)**: This function returns confidence score of the city name of the IP address.
- **getLatitude(<ip>)**: This function returns the latitude associated with the IP address.
- **getLongitude(<ip>)**: This function returns the longitude associated with the IP address.
- **getTimezone(<ip>)**: This function returns the timezone associated with the IP address.
- **getAccuracyRadius(<ip>)**: This function returns the accuracy radius associated with the IP address, 0 if unknown.
- **getAverageIncome(<ip>)**: This function returns the average income of the region associated with the IP address, 0 if unknown.
- **getMetroCode(<ip>)**: This function returns the metro code of the region associated with the IP address, 0 if unknown.
- **getPopulationDensity(<ip>)**: This function returns the population density associated with the IP address.
- **getPostalCode(<ip>)**: This function returns the postal code associated with the IP address.
- **getCoordPoint(<ip>)**: This function returns a point for use in GIS functions of the lat/long of associated with the IP address.
- **getASN(<ip>)**: This function returns the autonomous system of the IP address, "Unknown" if the IP is unknown or invalid.
- **getASNOrganization(<ip>)**: This function returns the autonomous system organization of the IP address, "Unknown" if the IP is unknown or invalid.
- **isEU(<ip>), isEuropeanUnion(<ip>)**: This function returns `true` if the ip address is located in the European Union, `false` if not.
- **isAnonymous(<ip>)**: This function returns `true` if the ip address is anonymous, `false` if not.
- **isAnonymousVPN(<ip>)**: This function returns `true` if the ip address is an anonymous virtual private network (VPN), `false` if not.
- **isHostingProvider(<ip>)**: This function returns `true` if the ip address is a hosting provider, `false` if not.
- **isPublicProxy(<ip>)**: This function returns `true` if the ip address is a public proxy, `false` if not.
- **isTORExitNode(<ip>)**: This function returns `true` if the ip address is a known TOR exit node, `false` if not.

<https://github.com/cgivre/drill-geoip-functions>

Who is Trying to Hack This Site?

```
SELECT request_receive_time, connection_client_host, request_useragent,  
getCountryName(connection_client_host ) as countryName,  
getCityName(connection_client_host ) as cityName  
FROM dfs.test.`hackers-access.httpd`  
WHERE request_firstline_uri = '/join_form'
```

Explore

.CSV

Clipboard

Search Results

connection_client_host	request_useragent	countryName	cityName
23.95.237.180	Mozilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20100101 Firefox/35.0	United States	Buffalo
23.95.237.180	Mozilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20100101 Firefox/35.0	United States	Buffalo
158.222.5.157	Mozilla/5.0 (Windows NT 6.3; WOW64; rv:34.0) Gecko/20100101 Firefox/34.0 AlexaToolbar/alxf-2.21	United States	Wilmington
158.222.5.157	Mozilla/5.0 (Windows NT 6.3; WOW64; rv:34.0) Gecko/20100101 Firefox/34.0 AlexaToolbar/alxf-2.21	United States	Wilmington
5.39.5.5	Mozilla/5.0 (Windows NT 5.1; rv:34.0) Gecko/20100101 Firefox/34.0	France	Unknown
180.180.64.16	Mozilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20100101 Firefox/35.0	Thailand	Pattaya
180.180.64.16	Mozilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20100101 Firefox/35.0	Thailand	Pattaya
89.42.237.71	Mozilla/5.0 (Windows NT 5.1; rv:34.0) Gecko/20100101 Firefox/34.0	Spain	Roldan
89.42.237.71	Mozilla/5.0 (Windows NT 5.1; rv:34.0) Gecko/20100101 Firefox/34.0	Spain	Roldan
216.158.199.158	Mozilla/5.0 (Windows NT 5.1; rv:34.0) Gecko/20100101 Firefox/34.0	United States	Unknown

Who is Trying to Hack This Site?

Access by Country ☆ Altered

30 rows

00:00:03.08




.json



.csv



Canada China Germany Thailand United States



Who is Trying to Hack This Site?

```
SELECT request_receive_time, connection_client_host, request_useragent,  
getCountryName(connection_client_host ) as countryName,  
getCityName(connection_client_host ) as cityName,  
getLatitude(connection_client_host ) as latitude,  
getLongitude(connection_client_host ) as longitude
```

request_useragent	countryName	cityName	latitude	longitude
zilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20100101 Firefox/35.0	United States	Buffalo	42.8864	-78.8781
zilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20100101 Firefox/35.0	United States	Buffalo	42.8864	-78.8781
zilla/5.0 (Windows NT 6.3; WOW64; rv:34.0) Gecko/20100101 Firefox/34.0 AlexaToolbar/alxf-2.21	United States	Wilmington	39.8188	-75.5064
zilla/5.0 (Windows NT 6.3; WOW64; rv:34.0) Gecko/20100101 Firefox/34.0 AlexaToolbar/alxf-2.21	United States	Wilmington	39.8188	-75.5064
zilla/5.0 (Windows NT 5.1; rv:34.0) Gecko/20100101 Firefox/34.0	France	Unknown	48.8582	2.3387000000000002
zilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20100101 Firefox/35.0	Thailand	Pattaya	13.05	100.9333
zilla/5.0 (Windows NT 5.1; rv:35.0) Gecko/20100101 Firefox/35.0	Thailand	Pattaya	13.05	100.9333
zilla/5.0 (Windows NT 5.1; rv:34.0) Gecko/20100101 Firefox/34.0	Spain	Roldan	37.798	-1.0097
zilla/5.0 (Windows NT 5.1; rv:34.0) Gecko/20100101 Firefox/34.0	Spain	Roldan	37.798	-1.0097
zilla/5.0 (Windows NT 5.1; rv:34.0) Gecko/20100101 Firefox/34.0	United States	Unknown	37.751	-97.822

What equipment are they
using?

What equipment are they using?

- Drill has support for multidimensional data structures including KV Pairs and Lists.
- There is a pre-existing UDF (<https://github.com/nielsbasjes/yauaa>) for Apache Drill which can parse User Agent Strings and get you a lot of useful information from the UA string.

Who is Trying to Hack This Site?


```
SELECT parse_user_agent(request_useragent) AS ua
FROM dfs.test.`hackers-access.httpd`
WHERE request_firstline_uri = '/join_form'
```

 Explore	 .CSV	 Clipboard	Search Results
ua			
{"DeviceClass":"Desktop","DeviceName":"Desktop","DeviceBrand":"Unknown","DeviceCpuBits":"32","OperatingSystemClass":"Desktop","OperatingSystemName":"Windows NT","OperatingSystemVersion":"XP","Operati			
{"DeviceClass":"Desktop","DeviceName":"Desktop","DeviceBrand":"Unknown","DeviceCpuBits":"32","OperatingSystemClass":"Desktop","OperatingSystemName":"Windows NT","OperatingSystemVersion":"XP","Operati			
{"DeviceClass":"Desktop","DeviceName":"Desktop","DeviceBrand":"Unknown","DeviceCpuBits":"64","OperatingSystemClass":"Desktop","OperatingSystemName":"Windows NT","OperatingSystemVersion":"8.1","Operati			
{"DeviceClass":"Desktop","DeviceName":"Desktop","DeviceBrand":"Unknown","DeviceCpuBits":"64","OperatingSystemClass":"Desktop","OperatingSystemName":"Windows NT","OperatingSystemVersion":"8.1","Operati			
{"DeviceClass":"Desktop","DeviceName":"Desktop","DeviceBrand":"Unknown","DeviceCpuBits":"32","OperatingSystemClass":"Desktop","OperatingSystemName":"Windows NT","OperatingSystemVersion":"XP","Operati			
{"DeviceClass":"Desktop","DeviceName":"Desktop","DeviceBrand":"Unknown","DeviceCpuBits":"32","OperatingSystemClass":"Desktop","OperatingSystemName":"Windows NT","OperatingSystemVersion":"XP","Operati			
{"DeviceClass":"Desktop","DeviceName":"Desktop","DeviceBrand":"Unknown","DeviceCpuBits":"32","OperatingSystemClass":"Desktop","OperatingSystemName":"Windows NT","OperatingSystemVersion":"XP","Operati			
{"DeviceClass":"Desktop","DeviceName":"Desktop","DeviceBrand":"Unknown","DeviceCpuBits":"32","OperatingSystemClass":"Desktop","OperatingSystemName":"Windows NT","OperatingSystemVersion":"XP","Operati			
{"DeviceClass":"Desktop","DeviceName":"Desktop","DeviceBrand":"Unknown","DeviceCpuBits":"32","OperatingSystemClass":"Desktop","OperatingSystemName":"Windows NT","OperatingSystemVersion":"XP","Operati			

Who is Trying to Hack This Site?

```
SELECT table1.ua.OperatingSystemNameVersion as os
table1.ua.AgentNameVersionMajor as browser
FROM
(
SELECT
parse_user_agent(request_useragent) AS ua
FROM dfs.test.`hackers-access.httpd`
WHERE request_firstline_uri = '/join_form'
) AS table1
```

Browser Breakdown ☆



Firefox 34	Clipboard
os	browser
Windows XP	Firefox 35
Windows XP	Firefox 35
Windows 8.1	AlexaToolbar alxf
Windows 8.1	AlexaToolbar alxf
Windows XP	Firefox 34
Windows XP	Firefox 35
Windows XP	Firefox 35
Windows XP	Firefox 34
Windows XP	Firefox 34
Windows XP	Firefox 34

Questions so far?

Let's have some fun with
PCAP

Let's have some fun with PCAP

- PCAP is short for **P**acket **C**apture and represent raw network traffic.
- Files are encoded in binary format, but various tools exist to analyze PCAP files or convert them into more easily accessible formats, such as JSON.

PCAP Analysis

arp-storm.pcap	
type	VARCHAR
network	INTEGER
timestamp	TIMESTAMP
timestamp_micro	BIGINT
src_ip	VARCHAR
dst_ip	VARCHAR
src_port	INTEGER
dst_port	INTEGER
src_mac_address	VARCHAR
dst_mac_address	VARCHAR
tcp_session	BIGINT
tcp_ack	INTEGER
tcp_flags	INTEGER
tcp_flags_ns	INTEGER
tcp_flags_cwr	INTEGER
tcp_flags_ece	INTEGER
tcp_flags_ece_ecn_capable	INTEGER
tcp_flags_ece_congestion_experienced	INTEGER
tcp_flags_urg	INTEGER
tcp_flags_ack	INTEGER
tcp_flags_psh	INTEGER
tcp_flags_rst	INTEGER
tcp_flags_syn	INTEGER
tcp_flags_fin	INTEGER
tcp_parsed_flags	VARCHAR
packet_length	INTEGER
data	VARCHAR

PCAP Analysis

- Drill has a comprehensive collection of networking functions to facilitate network analysis
- `is_private_ip()`, `is_valid_ip()`, `in_network()` and many others.
- `inet_aton()` and `inet_ntoa()` exist to facilitate sorting IP ranges.

PCAP Analysis


type	network	timestamp	timestamp_micro	src_ip	dst_ip	src_port	dst_port	src_mac_address	dst_mac_address
2005-03-30T08:47:50.501000									
UDP	1	2005-03-30T08:47:46.496000	1112172466496576	192.168.170.20	192.168.170.8	53	32795	00:C0:9F:32:41:8C	00:E0:18:B1:0C:AD
UDP	1	2005-03-30T08:47:50.501000	1112172470501268	192.168.170.8	192.168.170.20	32795	53	00:E0:18:B1:0C:AD	00:C0:9F:32:41:8C
UDP	1	2005-03-30T08:47:51.333000	1112172471333401	192.168.170.20	192.168.170.8	53	32795	00:C0:9F:32:41:8C	00:E0:18:B1:0C:AD
UDP	1	2005-03-30T08:47:59.313000	1112172479313231	192.168.170.8	192.168.170.20	32795	53	00:E0:18:B1:0C:AD	00:C0:9F:32:41:8C
UDP	1	2005-03-30T08:47:59.452000	1112172479452255	192.168.170.20	192.168.170.8	53	32795	00:C0:9F:32:41:8C	00:E0:18:B1:0C:AD
UDP	1	2005-03-30T08:48:07.320000	1112172487320873	192.168.170.8	192.168.170.20	32795	53	00:E0:18:B1:0C:AD	00:C0:9F:32:41:8C
UDP	1	2005-03-30T08:48:07.321000	1112172487321379	192.168.170.20	192.168.170.8	53	32795	00:C0:9F:32:41:8C	00:E0:18:B1:0C:AD
UDP	1	2005-03-30T08:49:18.685000	1112172558685951	192.168.170.8	192.168.170.20	32795	53	00:E0:18:B1:0C:AD	00:C0:9F:32:41:8C
UDP	1	2005-03-30T08:49:18.734000	1112172558734862	192.168.170.20	192.168.170.8	53	32795	00:C0:9F:32:41:8C	00:E0:18:B1:0C:AD
UDP	1	2005-03-30T08:49:35.461000	1112172575461181	192.168.170.8	192.168.170.20	32795	53	00:E0:18:B1:0C:AD	00:C0:9F:32:41:8C

Let's look for who is talking
to whom...


PCAP Analysis

```
SELECT DISTINCT src_ip, dst_ip
FROM dfs.test.`slowdownload.pcap`
WHERE src_ip is not null
```


src_ip	dst_ip
10.100.252.1	255.255.255.255
10.0.52.164	216.251.114.10
10.0.52.164 4.10	10.0.52.164
10.0.52.164	66.218.71.162
66.218.71.162	10.0.52.164
10.0.52.164	66.102.7.104
66.102.7.104	10.0.52.164
10.0.52.164	202.177.217.153
202.177.217.153	10.0.52.164
10.0.52.164	63.150.131.19
63.150.131.19	10.0.52.164
10.0.52.164	63.150.131.17



SynFlood Detection



SynFlood Detection



Drill can automate SYN Flood detection with a custom UDF.

<https://github.com/cgivre/drill-synflood-udf>

SYN Flood Detection

```
SELECT tcp_session
FROM dfs.test.`synscan.pcap`
GROUP BY tcp_session
HAVING is_syn_flood(tcp_session, tcp_flags_syn, tcp_flags_ack)
```

tcp_session
6346604732028469374
-9031405983396365775
7738739733723725373

SYN Flood Detection

```
SELECT *
FROM dfs.test.`synscan.pcap`
WHERE tcp_session IN
(
  SELECT tcp_session
  FROM dfs.test.`synscan.pcap`
  GROUP BY tcp_session
  HAVING is_syn_flood(tcp_session, tcp_flags_syn, tcp_flags_ack)
)
```

type	network	timestamp	timestamp_micro	src_ip	dst_ip	src_port	dst_port	src_mac_address	dst_mac_address	tcp_session
TCP	1	2010-07-04T20:24:16.276000	1278275056276783	172.16.0.8	64.13.134.52	36050	53	00:25:B3:BF:91:EE	00:26:0B:31:07:33	6346604732028469374
TCP	1	2010-07-04T20:24:16.338000	1278275056338667	64.13.134.52	172.16.0.8	53	36050	00:26:0B:31:07:33	00:25:B3:BF:91:EE	6346604732028469374
TCP	1	2010-07-04T20:24:16.403000	1278275056403870	172.16.0.8	64.13.134.52	36050	80	00:25:B3:BF:91:EE	00:26:0B:31:07:33	-9031405983396365775
TCP	1	2010-07-04T20:24:16.464000	1278275056464845	64.13.134.52	172.16.0.8	80	36050	00:26:0B:31:07:33	00:25:B3:BF:91:EE	-9031405983396365775
TCP	1	2010-07-04T20:24:17.678000	1278275057678354	172.16.0.8	64.13.134.52	36050	22	00:25:B3:BF:91:EE	00:26:0B:31:07:33	7738739733723725373
TCP	1	2010-07-04T20:24:17.740000	1278275057740531	64.13.134.52	172.16.0.8	22	36050	00:26:0B:31:07:33	00:25:B3:BF:91:EE	7738739733723725373
TCP	1	2010-07-04T20:24:19.338000	1278275059338245	64.13.134.52	172.16.0.8	53	36050	00:26:0B:31:07:33	00:25:B3:BF:91:EE	6346604732028469374
TCP	1	2010-07-04T20:24:19.462000	1278275059462133	64.13.134.52	172.16.0.8	80	36050	00:26:0B:31:07:33	00:25:B3:BF:91:EE	-9031405983396365775
TCP	1	2010-07-04T20:24:21.338000	1278275061338288	64.13.134.52	172.16.0.8	22	36050	00:26:0B:31:07:33	00:25:B3:BF:91:EE	7738739733723725373

Questions so far?

How about log files?


```
Dec 12 2018 06:50:25 sshd[36669]: Failed password for root from 61.160.251.136 port 1313 ssh2
Dec 12 2018 06:50:25 sshd[36669]: Failed password for root from 61.160.251.136 port 1313 ssh2
Dec 12 2018 06:50:24 sshd[36669]: Failed password for root from 61.160.251.136 port 1313 ssh2
Dec 12 2018 03:36:23 sshd[41875]: Failed password for root from 222.189.239.10 port 1350 ssh2
Dec 12 2018 03:36:22 sshd[41875]: Failed password for root from 222.189.239.10 port 1350 ssh2
Dec 12 2018 03:36:22 sshlockout[15383]: Locking out 222.189.239.10 after 15 invalid attempts
Dec 12 2018 03:36:22 sshd[41875]: Failed password for root from 222.189.239.10 port 1350 ssh2
Dec 12 2018 03:36:22 sshlockout[15383]: Locking out 222.189.239.10 after 15 invalid attempts
Dec 12 2018 03:36:22 sshd[42419]: Failed password for root from 222.189.239.10 port 2646 ssh2
```

eventDate	process_name	pid	message	src_ip
2018-12-12T11:50:25	sshd	36669	Failed password for root from 61.160.251.136 port 1313 ssh2	61.160.251.136
2018-12-12T11:50:25	sshd	36669	Failed password for root from 61.160.251.136 port 1313 ssh2	61.160.251.136
2018-12-12T11:50:24	sshd	36669	Failed password for root from 61.160.251.136 port 1313 ssh2	61.160.251.136
2018-12-12T08:36:23	sshd	41875	Failed password for root from 222.189.239.10 port 1350 ssh2	222.189.239.10
2018-12-12T08:36:22	sshd	41875	Failed password for root from 222.189.239.10 port 1350 ssh2	222.189.239.10
2018-12-12T08:36:22	sshlockout	15383	Locking out 222.189.239.10 after 15 invalid attempts	222.189.239.10
2018-12-12T08:36:22	sshd	41875	Failed password for root from 222.189.239.10 port 1350 ssh2	222.189.239.10
2018-12-12T08:36:22	sshlockout	15383	Locking out 222.189.239.10 after 15 invalid attempts	222.189.239.10
2018-12-12T08:36:22	sshd	42419	Failed password for root from 222.189.239.10 port 2646 ssh2	222.189.239.10
2018-12-12T08:36:22	sshlockout	15383	Locking out 222.189.239.10 after 15 invalid attempts	222.189.239.10
2018-12-12T08:36:22	sshd	42419	Failed password for root from 222.189.239.10 port 2646 ssh2	222.189.239.10

Drill can natively query
Syslog formatted data*

*** This feature will be available in Drill version 1.16**

Questions?



<https://amzn.to/2HDwE92>

Drill and DataDistillr

- Drill is a FOSS product. You can download and install at <https://drill.apache.org>.
- DataDistillr is a commercial SaaS product built on top of Drill.
- DataDistillr runs the bleeding edge version of Drill, so there are some features that are not yet available in the current stable version of Drill.

Advanced SQL Functionality

Advanced SQL Functionality

DataDistlr/Drill uses a SQL Dialect that is very similar to MySQL, so if you are comfortable with that, you'll find the learning curve to be minimal. However, there are a few key areas where it diverges significantly:

- Files
- Nested/Complex Data
- Remote Data
- Specialized Functionalities

Querying Files

- Querying files is fundamentally no different in DataDistillr than querying a database, however the FROM clause is a little different.
- When querying a database, typically the from clause will look like this:

```
FROM <catalog>.<database>.<table>
```

- When you query a file the components are a little different:

```
FROM `<filesystem>`.`<optional workspace>`.`<somefile.csv>`
```

- Note that Drill uses backticks for identifiers.

Querying Files

- The workspace is an optional shortcut to a file path.
- You can query directories, as well as glob files.

Examples:

```
FROM dfs.`/path/to/my/files/file.json`
```

```
FROM dropbox.`my_logs/*.log`
```

```
FROM box.`my_data/`
```

Querying Files

- Almost every file format has configurable options which can be set globally in the Drill config or at query time by using the table function.
- The options are different for every format, however documentation is located in the Drill docs: (<https://drill.apache.org/docs/querying-a-file-system-introduction/>) or in Drill's github repository: (<https://github.com/apache/drill/tree/master/contrib>)
- For example, Excel and MS Access files contain more than one table. You can choose which table to query with the following syntax:

```
SELECT <fields>  
FROM table(dfs.`test_data.xlsx` (type => 'excel', sheetName => 'secondSheet'))
```

Querying Files

- Try it yourself! Take a few minutes to try querying some of the files in the demo data folder.

Nested/Complex Data

- One of the fundamental principles of database design is normalization, which states that each table "cell" should be atomic, and thus have only one piece of data in it.
- SQL was originally designed with this in mind.
- Modern SQL systems have extensions to allow users to query nested & complex data. Unfortunately, the syntax is not standardized across platforms.
- Drill supports two data structures: lists & dictionaries. AKA arrays and maps.

Nested/Complex Data: Lists

If you have data that contains a list, you can access individual elements using bracket notation:

```
SELECT my_list[0] as field_1, my_list[1] AS field_2
```

If you have nested lists, the syntax is the same:

```
SELECT my_nested_list[1][0], my_nested_list[1][2]
```

Nested/Complex Data: Lists

You can also transform lists or repeated maps using the FLATTEN function.

For instance, if you have a list [1,2,3] called my_list, you could write a query:

```
SELECT FLATTEN(my_list) ...
```

This will create a row for every entry in the list.

Nested/Complex Data

DataDistillr also supports Key/Value AKA maps AKA associative arrays.

DataDistillr supports two syntax for accessing map elements:

```
SELECT my_map.field
```

OR

```
SELECT my_map['field']
```

My recommendation is to use the bracket notation rather than the dotted notation.

There is also a function `KVGEN()`, which when used with `FLATTEN()` can be used to transform maps into columns.

Nested/Complex Data

- Using the `http-pcap.json`, try out some queries to extract nested fields. Some questions include:
- What are the different source IP addresses?
- How many different MAC addresses are there?

Specialized Functionalities

- DataDistillr has many specialized functions for data cleaning and manipulation, not all are OSS.
- You've already seen the IP geolocation functions, there are also functions for working with phone numbers, street addresses and others

Specialized Functions: URLs

- `parse_url(<url>)`: This function accepts a URL as an argument and returns a map of the URL's protocol, authority, host, and path.
- `parse_query(<query_string>)`: This function accepts a query string and returns a key/value pairing of the variables submitted in the request.
- Try it out using the `url_log.csv` file. See if you can extract all the query parameters as kv pairs

Specialized Functions: DNS

- `getHostName(<IP address>)`: Returns the host name associated with an IP address.
- `getHostAddress(<host>)`: Returns an IP address associated with a host name.
- `dnsLookup(<host>, [<Resolver>])`: Performs a DNS lookup on a given host. You can optionally provide a resolver. Possible resolver values are: `cloudflare`, `cloudflare_secondary`, `google`, `google_secondary`, `verisign`, `verisign_secondary`, `yandex`, `yandex_secondary`.
- `whois(<host>, [<Resolver>])`: Performs a whois lookup on the given host name. You can optionally provide a resolver URL. Note that not all providers allow bulk automated whois lookups, so please follow the terms of service for your provider.

Querying Remote Data

- DataDistillr can query remote data VIA API calls.
- DataDistillr will support virtually any API as long as the API returns JSON, CSV or XML.
- DataDistillr supports many different options for APIs including pagination, OAuth 2.0 and more.
- When properly configured, DataDistillr will push down filters to the downstream API.

Thank You!!

Charles S. Givre, CISSP
charles.givre@gtkcyber.com
www.thedataist.com